
ENPM 667 - Project 1 - Implementation of Journal Paper :

Low-level autonomous control and tracking of quadrotor using reinforcement learning

By:-

Anubhav Paras - 116905909

Preyash Rajeshkumar Parikh - 117303698

Contents

- Overview
 - Background
 - Introduction
 - Understanding
 - Algorithm - Training and Result
 - Experimental Results- Training and Result
 - Comparison of Results
 - Quadrotor Dynamics
 - References
-

Overview of Project - Concept



The concept is about implementing a low level flight controller which uses Proximal Policy Optimization where we require the controller not to deviate vastly from the previous policy while aiming to achieve a stable policy. This has to be carried out while maximizing the expected return. Hence, this is a smooth learning process.

Background

- **UAV and Applications !**

→ Patrolling, Aerial monitoring and others

- **Issues with Traditional and Optimization Methods.**

→ Traditional methods are required to have parameter adjustments

→ In Optimization algorithms, quadrotor is subjected to external disturbance.

- **What we proposed ?**

→ Proximal Policy Optimization



Introduction

- **Flight control systems of quadrotors**
 - Outer Loop - Higher- Level Planner
 - Inner Loop - Lower Level Controller
 - **Why reinforcement learning based strategy ?**
 - To have a no trial-and-error tuning of parameters.
 - **Why Model Free methods ?**
 - Since domain knowledge is not required.
-

UNDERSTANDING


→ **Markov Decision Process**

→ **Policy Gradient**

→ **Proximal Policy Optimization**

Markov Decision Process

The Markov property states that, “ The future is independent of the past given the present.” Once we know the current state, the past information encountered may be thrown away, and that state is a sufficient statistic that gives us the same view of the future as if we have all the past information.




An MDP is a tuple (S, A, P, R, γ)

$$\mathcal{P}_{ss'}^a = \mathbb{P}[S_{t+1} = s' \mid S_t = s, A_t = a]$$

Policy Gradient

Policy gradient methods are a type of reinforcement learning techniques that rely upon optimizing parametrized policies with respect to the expected return.

Policy Gradient



$$p_{\theta}(s_1, a_1, \dots, s_T, a_T) = p(s_1) \prod_{t=1}^T \pi_{\theta}(a_t | s_t) p(s_{t+1} | s_t, a_t)$$

$\pi_{\theta}(\tau)$

$$\sum_{s \in \mathcal{S}, a \in \mathcal{A}} \rho^{\pi}(s) Q^{\pi}(s, a) \frac{\partial \pi(a|s)}{\partial \theta}$$

$$= \sum_{s \in \mathcal{S}, a \in \mathcal{A}} \rho^{\pi}(s) (V^{\pi}(s) + A^{\pi}(s, a)) \frac{\partial \pi(a|s)}{\partial \theta}$$


$$= \sum_{s \in \mathcal{S}, a \in \mathcal{A}} \rho^{\pi}(s) A^{\pi}(s, a) \frac{\partial \pi(a|s)}{\partial \theta},$$

Proximal Policy Optimization

→ **Problem with Policy Gradient** is that the training process is too slow and too much variability in training.

→ **PPO improves** the stability of the Actor training by limiting the policy update at each training step.

→ **Clipped Surrogate Function**



QUADROTOR DYNAMICS

- Takes gravity and the forces generated by motors.
- Six- Degree rigid body with four motor thrust forces.
- Torque generated is

$$\tau = J \begin{bmatrix} \frac{1}{\sqrt{2}} L(-F_1 + F_2 + F_3 - F_4) \\ \frac{1}{\sqrt{2}} L(F_1 - F_2 + F_3 - F_4) \\ K_m L(-F_1 + F_2 - F_3 + F_4) \end{bmatrix}$$

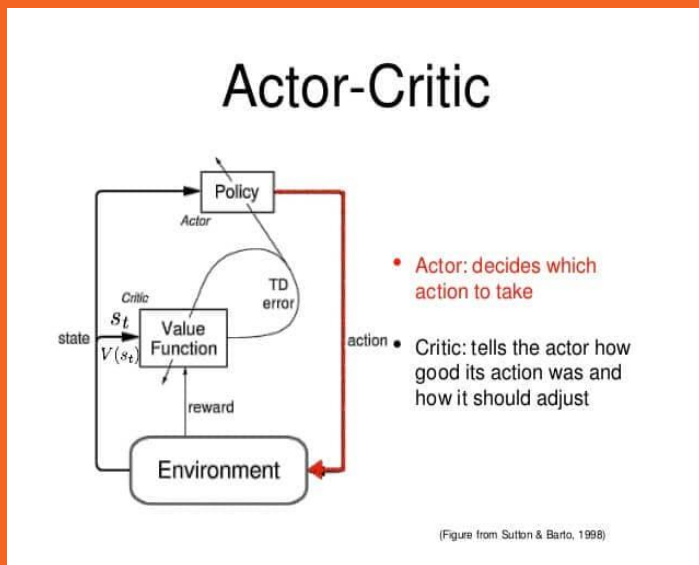
- Forced produce by motor spinning, $F_i = K_f \omega_i^2, \quad i = 1, 2, 3, 4$

Where, ω_i is the speed of the motor and

K_m is coefficient of the generated torque.

ALGORITHM - TRAINING

→ Policy Gradient based Actor- Critic architecture.



Value:

$$V^\pi(s_t) = \mathbb{E} \left[r_t + \gamma V^\pi(s_{t+1}) \mid a_t \sim \pi(a|s_t) \right]$$

Actor:

$$\theta_\pi \leftarrow \theta_\pi + \alpha \frac{1}{|B|} \sum_{(s_t, a_t, r_t, s_{t+1}) \in B} \frac{Q^\pi(s_t, a_t)}{\pi(a_t|s_t; \theta_\pi)} \frac{\partial \pi(a|s; \theta_\pi)}{\partial \theta_\pi},$$

Algorithm- Training

- The learning algorithm comprises of three main components:
 - Actor (Policy Function):
 - This is a simple neural network to approximate the parameterized policy.
 - Critic (Value Function):
 - Neural network to approximate the value function that returns the expected rewards given a state.
 - PPO (Proximal Policy Optimization) Agent: This is the heart of the algorithm that is responsible for maximizing the expected sum of rewards by imposing a constraint on the improved policies.

-

PPO Algorithm -

Algorithm 1 PPO, Actor-Critic Style

```
for iteration=1,2,... do
  for actor=1,2,...,N do
    Run policy  $\pi_{\theta_{\text{old}}}$  in environment for  $T$  timesteps
    Compute advantage estimates  $\hat{A}_1, \dots, \hat{A}_T$ 
  end for
  Optimize surrogate  $L$  wrt  $\theta$ , with  $K$  epochs and minibatch size  $M \leq NT$ 
   $\theta_{\text{old}} \leftarrow \theta$ 
end for
```

Tasks to train

- 1) Hovering:
 - a) The quadrotor is expected to hover over a specific target point: (0,0,10)
 - b) For this task the position error (x,y,z) and the orientation error(roll, pitch, yaw) should be minimum.
 - 2) Trajectory following:
 - a) The quadrotor is expected to circle around a fixed point in the x-y plane.
 - b) To achieve this it has to maintain a constant distance with a point (should be **within certain threshold** of the circle boundary) while also maintaining a constant desired velocity.
-

Algorithm- Training

For each EPISODE (E):

- 1) Collect a set of N (PPO steps) transitions (T) where a transition is defined by:
 - $T[i] = \langle \text{state}_i, \text{action}_i, \text{reward}_i, \text{curr_policy_distribution} \rangle$
 - i) Select an $\langle \text{action} \rangle$ from the **current policy** (actor)
 - ii) Execute the $\langle \text{action} \rangle$ to get the $\langle \text{next_state} \rangle$ and $\langle \text{reward} \rangle$
 - iii) $T[i] = \langle \text{state}, \text{action}, \text{reward}, \text{curr_policy_distribution} \rangle$
 - 2) Calculate returns using generalized advantage algorithm and the advantage $A = Q(s,a) - V(s)$: Advantage tells how good a particular action is as compared to an average action taken in some state.
 - 3) Use this batch of transitions and the advantages to train the actor-critic network to **update the policy**.
-

ALGORITHM - TRAINING

Training the actor-critic network:

- 1) In PPO to train the actor need to optimize the loss function given by:

$$L^{CLIP}(Q) = \hat{E}_t[\min(r_t(Q)\hat{A}_t, \text{clip}(r_t(Q), 1-\epsilon, 1+\epsilon)\hat{A}_t)]$$

- 2) Here we clip the ratio between the old and new policy by some factor, here $(1-0.2=0.8)$ to $(1+0.2=1.2)$.
- 3) For our implementation we modified the loss function as per the paper.

$$\begin{aligned} \text{maximize } L_{\text{policy}} &= \sum_{(s,a) \in T} \min \left[\left(\frac{\pi(a|s)}{\mu(a|s)} - 1 \right) A^{\text{trace}}, \epsilon |A^{\text{trace}}| \right] \\ \text{minimize } L_{\text{value}} &= \frac{1}{|T|} \sum_{(s,a) \in T} (V(s) - V^{\text{trace}})^2. \end{aligned}$$

ALGORITHM - TRAINING

Training the actor-critic network:

4) For each EPOCH while training:

- Sample the mini-batch from the batch of transitions (T)
- Feed forward the network to get action and policy distribution (from the actor), values (from the critic)
- Calculate the loss:
 - Surrogate loss for the actor = $\text{Adv} * \min [\exp \{\log(\text{new_distribution}) - \log(\text{old_distribution})\}, 1]$
 - MSE loss for the Critic = $(\text{returns} - V(s))^2$
- Backpropagate the total loss through the network using SGD

5) Repeat till convergence (ie old Policy and new policy do not diverge much)

ALGORITHM - TRAINING: Rewards

HOVERING:

- 1) For the task of reaching a particular target and hovering over there:
- $\text{reward} = w1 * (\text{position error}) + w2 * (\text{orientation error}) + w3 * (\text{action})$
 - 2) If the quadrotor reached the defined boundary limits a PENALTY was imposed.
 - 3) If the quadrotor reached within some threshold of the target position, a BONUS was given, where,
 $\text{bonus_reward} = \text{BONUS} * (\text{num_done}) * \text{reduction_factor}$
where num_done = number of times it reached within the threshold
-

ALGORITHM - TRAINING: Rewards

TRAJECTORY FOLLOWING:

- 1) For the task of reaching a particular target and hovering over there:
- reward = $1. w_1 * (\text{distance error}) + w_2 * (\text{velocity error}) + w_3 * (\text{action})$
 - 2) $\text{distance}_{\text{err}} = \text{desired_radius} - \sqrt{(x - x_{\text{center}})^2 - (y - y_{\text{center}})^2}$
 - 3) $\text{velocity}_{\text{err}} = x * y_{\text{vel}} - y * x_{\text{vel}} - \text{desired_radius} * \text{desired_vel}$
 - 4) If the quadrotor reached the defined boundary limits a PENALTY was imposed.
 - 5) If the quadrotor reached within some threshold of the target position, a BONUS was given, where,
bonus_reward = BONUS * (num_done) * reduction_factor
where **num_done = number of times it reached within the threshold**
-

ALGORITHM - TRAINING: Parameters

Quadrotor Params:

Mass	- 0.665kg
Arm Length	- 0.105m
I_{xx}	- 0.0023 kg-m ²
I_{yy}	- 0.0025 kg-m ²
I_{zz}	- 0.0037 kg-m ² m

Network Params

Actor network	- 32x32x4
Critic network	-128x128x1
Discount factor (gamma)	- 0.99
Learning rate (alpha)	-1e-4
PPO Steps	-350
Mini-Batch size	-64
Training epochs	-10/20

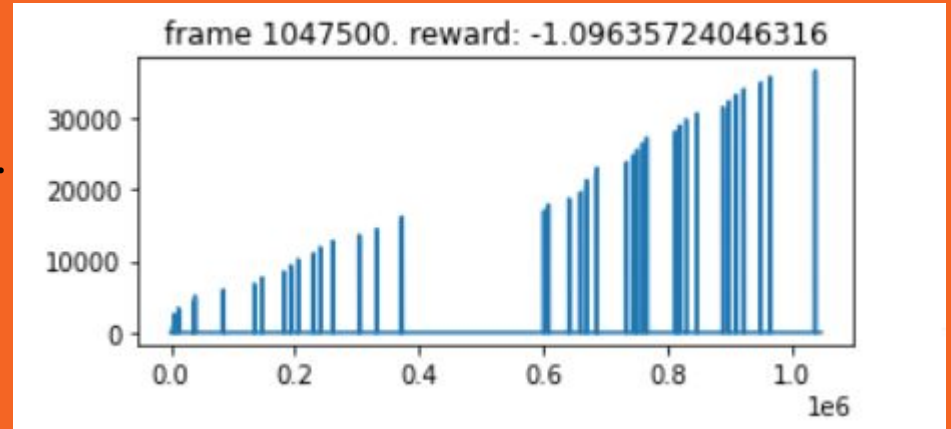
Input = state of the quadrotor = <position, angles, velocity, angular velocity>

Input state vector size = 1x12

RESULTS - TRAINING

Transition vs Rewards:


- We can observe that the rewards that the agent received were more in number as the training progressed.



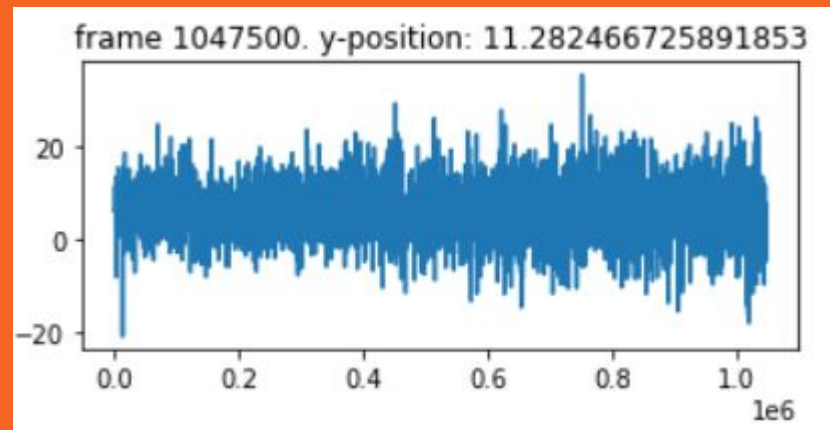

RESULTS - TRAINING

Transition vs Running Rewards:

- We observe that the total rewards started increasing as the training progressed.
- This can be attributed to the fact that more bonus rewards were obtained towards the end of training.





Results - Hovering Task




X and Y position were oscillating around 0 during training.

Results - Hovering Task



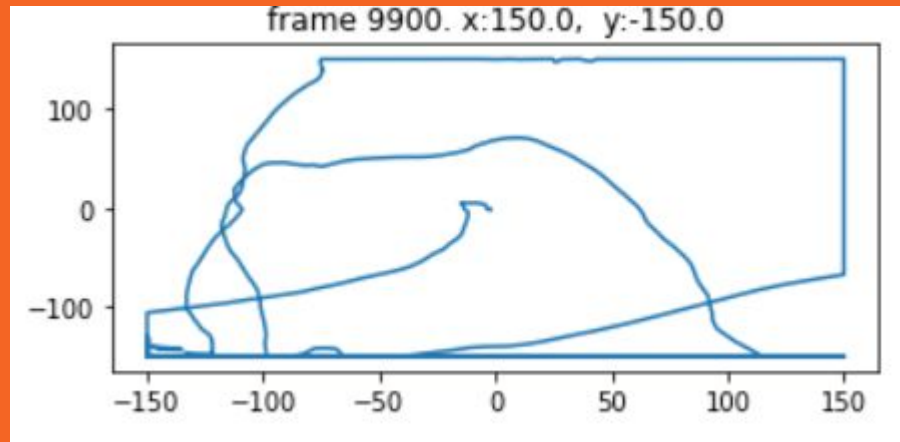
X and Y position oscillating around 0 during testing before converging.

Results - Hovering Task



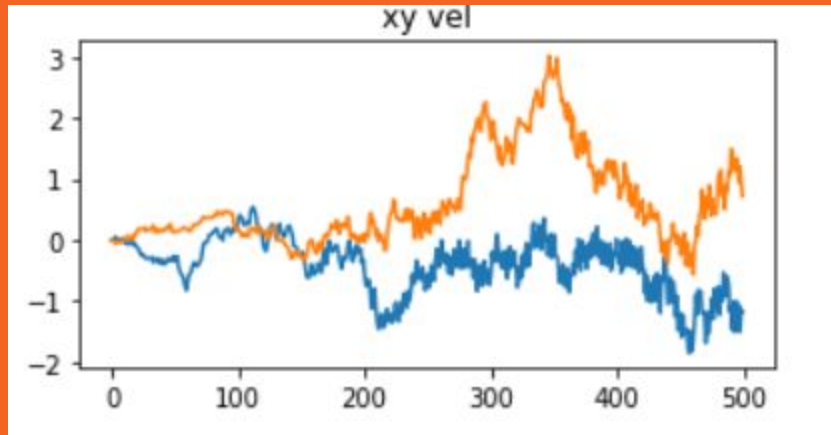
Results - Trajectory Following Task

The quadrotor encircle around the center of the circle.

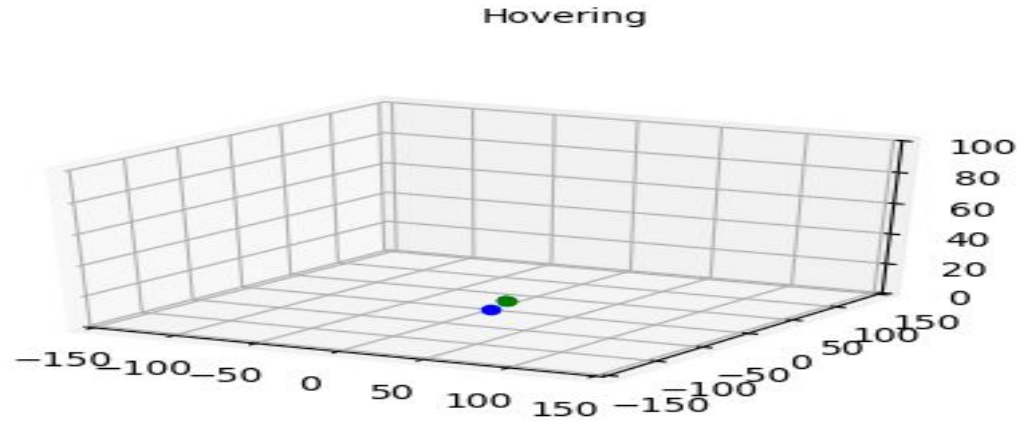


Results - Trajectory Following Task

X and Y component of the of the velocity were moving in a sinusoidal motion so as to maintain a constant desired velocity to move around in a circle.



Results - Trajectory Following Task



Results

- 1) For the hovering task training was done for $\sim 1\text{M}$ transitions (10K episodes of 350 PPO steps each).
 - 2) For the trajectory following task training was done for $\sim 0.3\text{M}$ transitions (1K episodes of 300 PPO steps each)
 - 3) Quadrotor was made to reach a point (0,0,0) and stay over there.
 - 4) It was following a near-to-circular trajectory.
-

Conclusion

- We tried to adhere to the functions, parameters and the algorithm as much as we can.
- The task results were satisfactory but can be improvised:
 - with more training episodes,
 - using multiple actors to provide more sample data at a time, i.e. using more diverse experience to improve the policy
 - modification to the loss function (calculating advantages using GAE - generalized advantage estimation),
 - tuning of few hyperparameters
- The training process took a lot of time after the implementation.
- To make it learn properly we had to run around 1 to 1.5M (~10M recommended in the paper) iterations and repeated that with each major or minor change in the code or the parameters to get better results.

-

References

- - 1) <https://towardsdatascience.com/reinforcement-learning-demystified-markov-decision-processes-part-1-bf00dda41690>
 - 2) <https://openai.com/blog/openai-baselines-ppo/>
 - 3) <https://jonathan-hui.medium.com/rl-proximal-policy-optimization-ppo-explained-77f014ec3f12>
 - 4) <https://stackoverflow.com/questions/46422845/what-is-the-way-to-understand-proximal-policy-optimization-algorithm-in-rl>
 - 5) <https://towardsdatascience.com/understanding-actor-critic-methods-931b97b6df3f>
 - 6) <https://www.cse.unsw.edu.au/~cs9417ml/RL1/tdlearning.html>
-

Deliverables

→ **Hover Code -**

https://colab.research.google.com/drive/15vZwsICSzSOU_38TRTV2KuEi7A5ZIt25?authuser=3#scrollTo=fCINsLY6DFvI

→ **Trajectory Code -**

https://colab.research.google.com/drive/1988BaVXe4V81RoQO3oeOvO7hvUTCN6BS?authuser=3#scrollTo=kmx_8l29LWJK

→ **Code -** https://github.com/anubhavparas/quadrotor_control_ppo

→ **Presentation**

→ **Report**

→ **Slides**

→ **Simulator GIF**

→ **Readme**
